

TensorFlow

for People Who Want to Use

TensorFlow

Overview

Install TensorFlow

TensorFlow Introduction

Linear Regression from scratch

Linear Regression, the easy way

Using Queues & Checkpoints

MNIST!

DeepDream, maybe?

Install TensorFlow

Follow the instructions at

<https://github.com/martinwicke/tensorflow-tutorial>

A multidimensional array.



TensorFlow

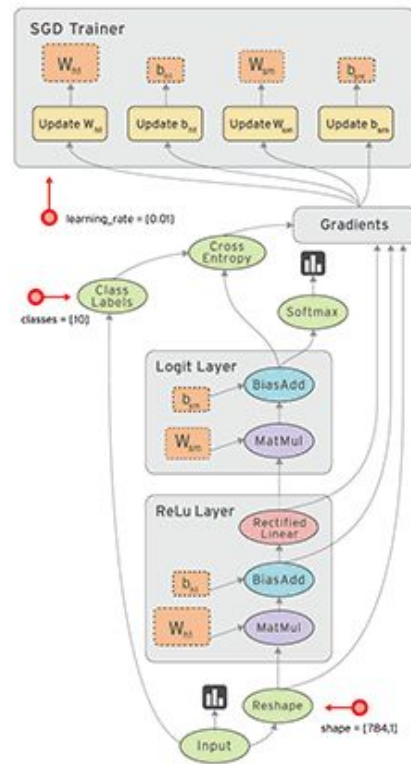


A graph of operations.

Flow

Computation is defined as a directed acyclic graph (DAG) to optimize an objective function

- Graph is defined in high-level language (Python)
- Graph is compiled and optimized
- Graph is executed (in parts or fully) on available low level devices (CPU, GPU)
- Data (tensors) flow through the graph
- TensorFlow can compute gradients automatically




Build a graph; then run it.

TensorFlow separates computation graph construction from execution.

...

```
c = tf.add(a, b)
```

```
session = tf.Session()  
numpy_c = c.eval(session)
```




What's in a Graph?

Edges: Tensors, Nodes: Ops

All nodes are Ops:

- Constants
- Variables
- Computation
- Debug code (Print, Assert)
- Control Flow



Variables

Some ops in a TensorFlow graph are stateful: (mainly) Variables

- Can be assigned to
- Must be initialized
- It is easy to create race conditions
 - Welcome to concurrent programming
 - Races are mostly harmless in stochastic data-parallel algorithms

Shape Inference

The data type of a Tensor is fixed during construction, its shape is not

```
a = tf.decode_png(bytes, channels=3,  
                  dtype=tf.uint8)
```

```
a.get_shape() ⇒ [None, None, 3]
```

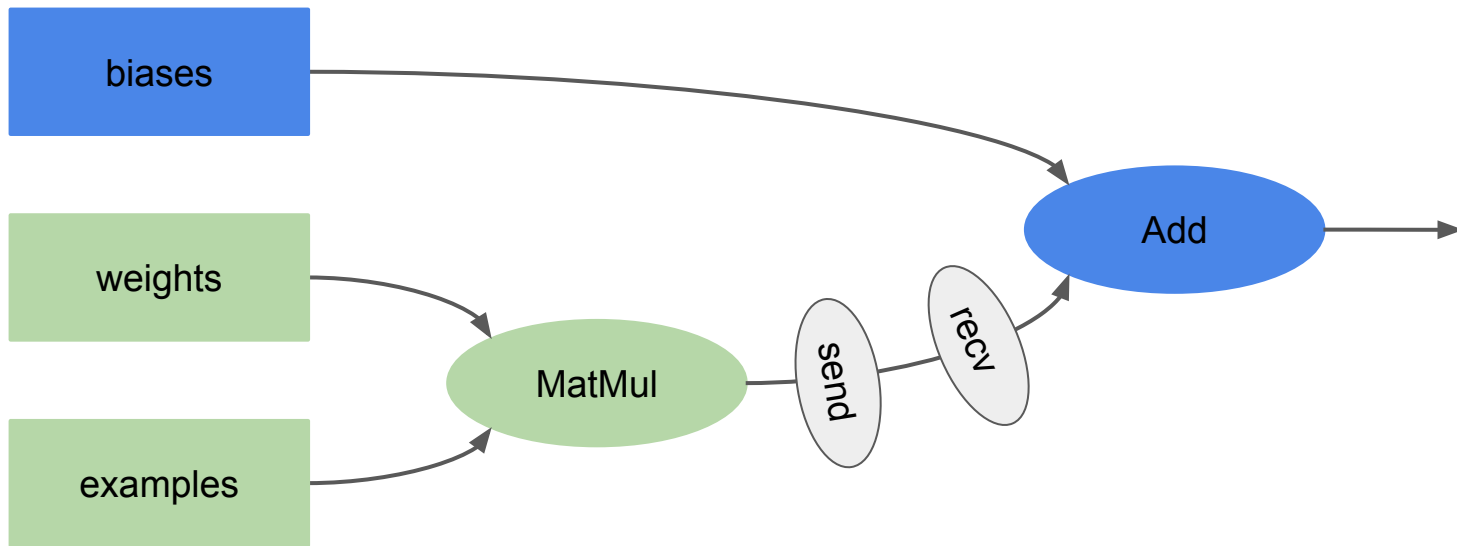
Shape inference propagates shapes as much as possible

```
b = tf.tile(a, [2, 2, 3])
```

```
b.get_shape() ⇒ [None, None, 9]
```

But, why?

Graphs can be processed, compiled, remotely executed, assigned to devices.



Automatic differentiation

Graph computing gradients can be computed automatically

Every Op has a corresponding gradient Op computing partial derivatives

TensorFlow knows the chain rule

You specify the forward computation, `tf.gradients` adds gradient Ops

Graphs can be explicit-ish

You can have several independent graphs at the same time

```
with tf.Graph().as_default():  
    a = tf.constant(1)  
    b = tf.constant(2)  
    c = tf.add(a, b)
```

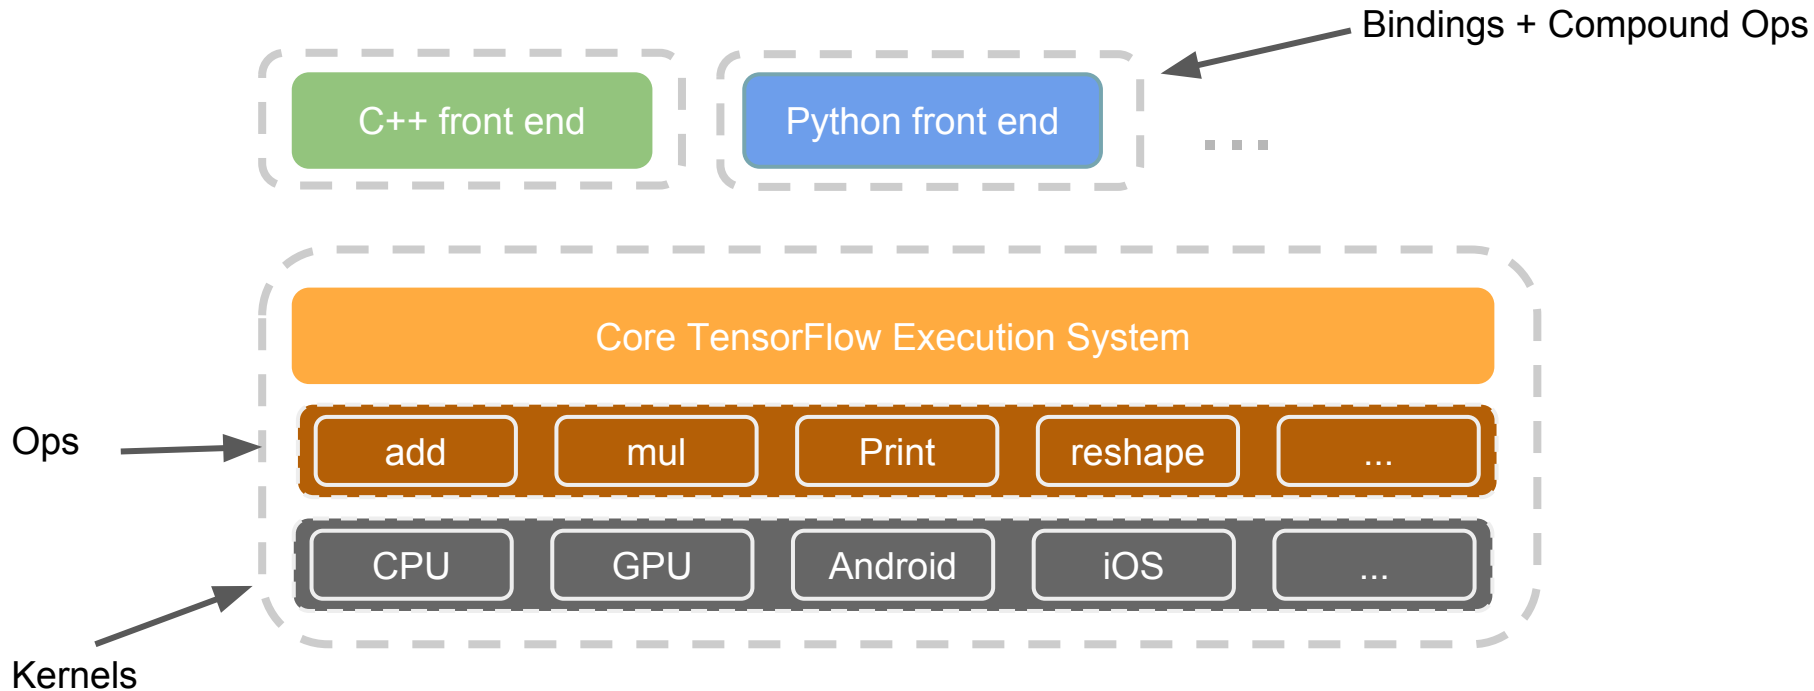
```
with tf.Graph().as_default():  
    a = tf.constant(10)  
    b = tf.add(a, c)    # Error!
```

Building graphs looks *mostly* like numpy

With special functions for deep learning

Numpy	TensorFlow
<code>add</code> <code>mul</code> <code>matmul</code> <code>...</code> <code>sum</code> <code>...</code>	<code>add</code> <code>mul</code> <code>matmul</code> <code>...</code> <code>reduce_sum</code> <code>...</code> <code>sigmoid</code> <code>relu</code> <code>...</code>

TensorFlow Architecture



Extending TensorFlow

Ops are small, but easy to combine into bigger pieces

- Writing new algorithms requires no knowledge about TensorFlow internals

```
def my_algorithm(input, depth):  
    output = input  
    for i in xrange(depth):  
        output = tf.contrib.layers.relu(output, 200)  
    return output
```

Extending TensorFlow

Ops are small, but easy to combine into bigger pieces

- Writing new algorithms requires no knowledge about TensorFlow internals
- Writing new compound Ops requires no knowledge about lower levels

```
def my_op(t, min, max):  
    t_min = math_ops.minimum(t, max)  
    t_max = math_ops.maximum(t_min, min)  
    return t_max
```


TensorFlow is a RISC architecture

Ops are small, but easy to combine into bigger pieces

- Writing new algorithms requires no knowledge about TensorFlow internals
- Writing new compound Ops requires no knowledge about lower levels

This flexibility makes TensorFlow ideal for Research

Let's write some code.

Linear Regression from Scratch


result

input

$$y = Wx + b$$

parameter

The diagram shows the equation $y = Wx + b$ on a white background. Three gray arrows point to the components of the equation: one from the word 'result' to the variable y , one from the word 'input' to the variable x , and one from the word 'parameter' to the term $Wx + b$.



$y = Wx + b$ in TensorFlow

```
import tensorflow as tf

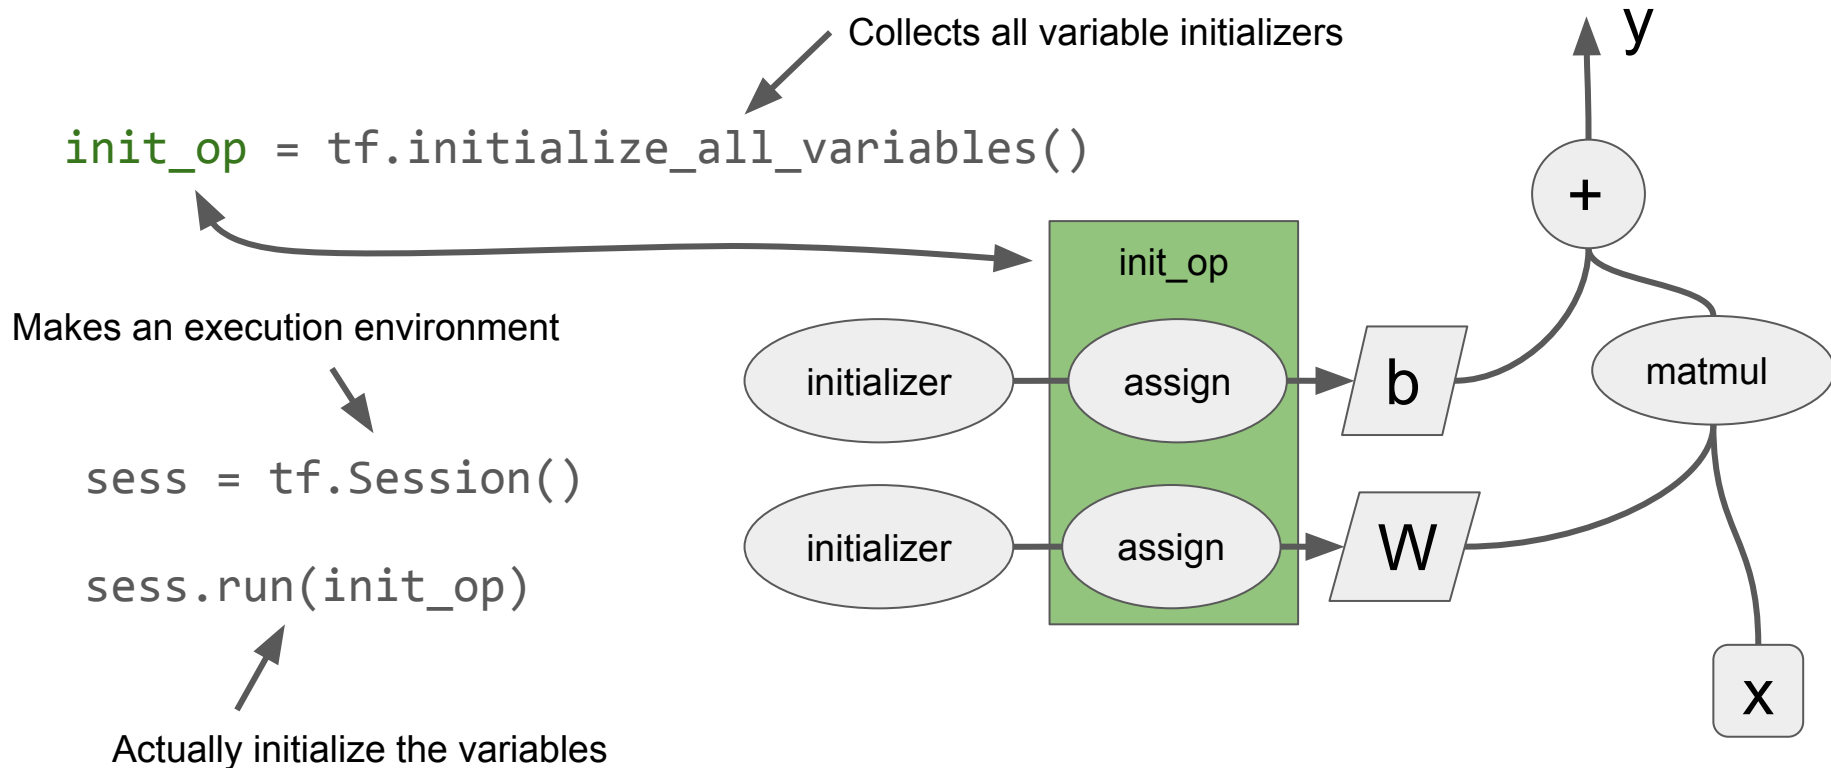
x = tf.placeholder(shape=[2,1], dtype=tf.float32, name="x")

W = tf.get_variable(shape=[1,2], name="W")

b = tf.get_variable(shape=[1], name="b")

y = tf.matmul(W, x) + b
```

Variables Must be Initialized




Running the Computation

```
x_in = [[3], [4]]
```

```
sess.run(y, feed_dict={x: x_in})
```

Only what's used to compute a fetch will be evaluated

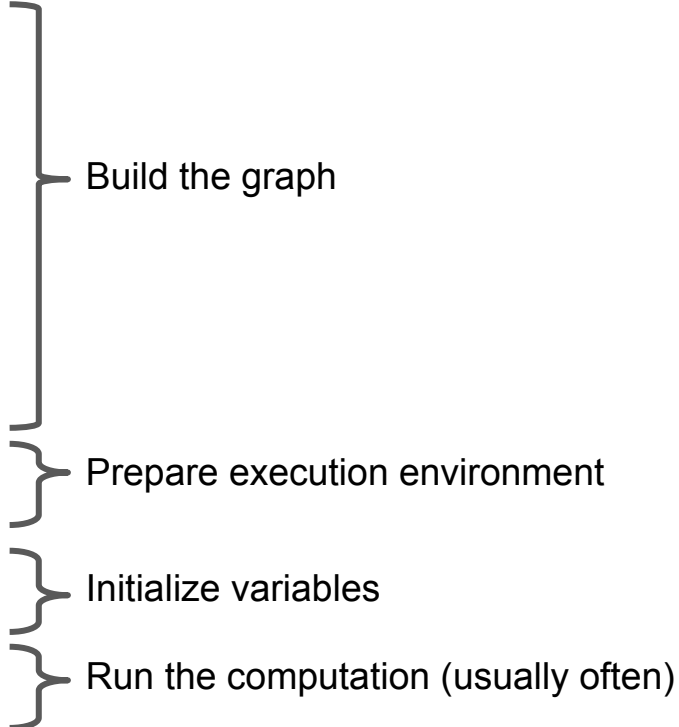
All Tensors can be fed, but all placeholders must be fed



The full program

```
import tensorflow as tf
x = tf.placeholder(shape=[2,1],
                   dtype=tf.float32,
                   name="x")
W = tf.get_variable(shape=[1,2], name="W")
b = tf.get_variable(shape=[1], name="b")
y = tf.matmul(W, x) + b

with tf.Session() as sess:
    sess.run(tf.initialize_all_variables())
    print sess.run(y, feed_dict={x: x_in})
```



The code is annotated with curly braces on the right side, grouping it into five stages:

- Build the graph:** This stage includes the definition of the placeholder `x`, the variables `W` and `b`, and the computation `y = tf.matmul(W, x) + b`.
- Prepare execution environment:** This stage includes the creation of the `tf.Session()` object.
- Initialize variables:** This stage includes the call to `sess.run(tf.initialize_all_variables())`.
- Run the computation (usually often):** This stage includes the call to `print sess.run(y, feed_dict={x: x_in})`.

Exercise: Define a Loss

Given \mathbf{x} , y_{label} , compute a loss, for instance:

$$L = (y - y_{\text{label}})^2$$

Hint, numpy's `sum` is called `reduce_sum`.

Solution

```
y_label = tf.placeholder(shape=[1,1], dtype=float32,  
                           name="y_label")  
  
diff = y - y_label  
L = tf.reduce_sum(diff * diff)
```

Evaluation

```
eval_data = np.loadtxt(open("eval_data.csv","rb"), delimiter=",")  
  
acc = 0.  
  
for x1, x2, y_in in eval_data:  
    acc += sess.run(L, feed_dict={x: [[x1],[x2]], y_label: y_in})  
  
print acc/len(eval_data)
```

Training

Feed $(\mathbf{x}, y_{\text{label}})$ pairs and adjust \mathbf{W} and \mathbf{b} to decrease the loss.

$$\mathbf{W} \leftarrow \mathbf{W} - \eta (dL/d\mathbf{W})$$

$$\mathbf{b} \leftarrow \mathbf{b} - \eta (dL/d\mathbf{b})$$

`tf.gradients(L, [W, b])` computes gradients of L .

`tf.GradientDescentOptimizer` creates Ops that perform the update step.

Training

```
L = ...
```

```
train_op = tf.train.GradientDescentOptimizer(learning_rate=0.01)  
            .minimize(L)
```

```
data = numpy.loadtxt(open("training_data.csv","rb"), delimiter=",")
```

```
for x1, x2, y_in in data:
```

```
    sess.run(train_op, feed_dict={x: [[x1],[x2]], y_label: y_in})
```

That seems complicated...

```
import tensorflow as tf
```

```
R = tf.contrib.learn.LinearRegressor(feature_columns=[  
    tf.contrib.layers.real_valued_column('', dimension=2)])
```

```
R.fit(x=data[:,0:2], y=data[:,2:3], batch_size=100, max_steps=100)
```

```
R.evaluate(x=eval_data[:,0:2], y=eval_data[:,2:3])
```

```
R.predict(x=np.asarray([1.5, 3.4]))
```

Exercise: Use a DNN

Hints: Use either

- Use `DNNRegressor`, or
- Start from the “From Scratch” version and use `tf.contrib.layers.relu`

Solution

```
import tensorflow as tf
```

```
D = tf.contrib.learn.DNNRegressor(feature_columns=[  
    tf.contrib.layers.real_valued_column('', dimension=2)],  
    hidden_units=[10,]*9)
```

Reading Data From Files

TensorFlow has text and binary file readers

`WholeFileReader`, `TestLineReader`, `FixedLengthRecordReader`, ...

As well as decoders

`decode_png`, `decode_jpg`, `decode_gif`, `decode_csv`, ...

The file readers need a `Queue` as input.

Queues

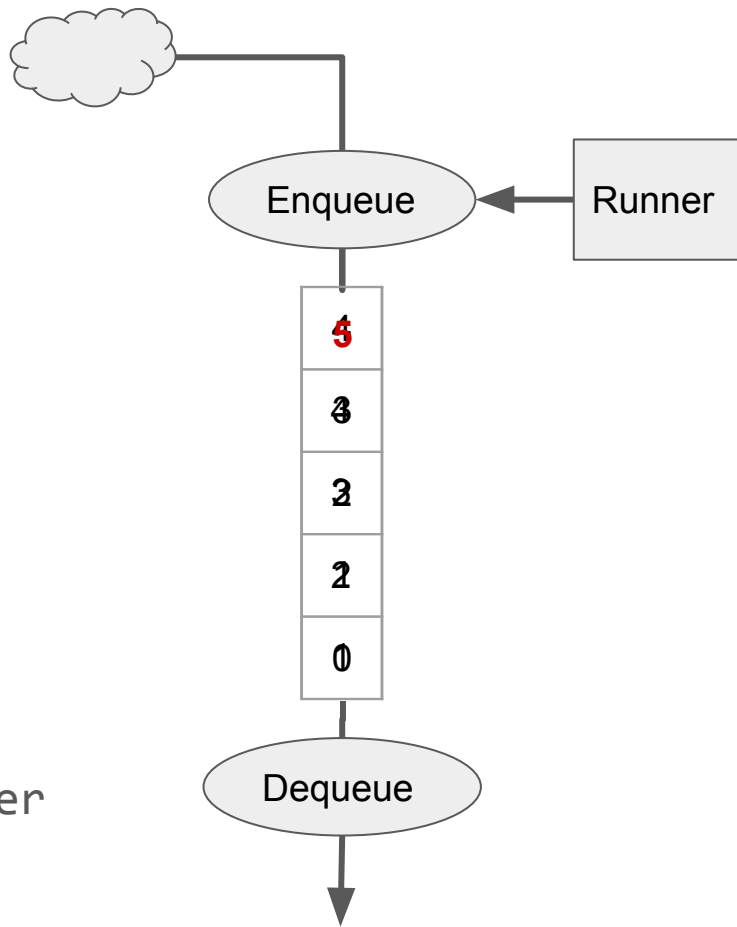
An Object with Enqueue, Dequeue Ops.

Can be asynchronously refilled by QueueRunners.

Internal state must be initialized!

State is not cleanly restored from Checkpoints!

Most common: `tf.train.string_input_producer`



Reading our data from file

```
filename_queue = tf.train.string_input_producer(["training_data.csv"], num_epochs=1)

reader = tf.TextLineReader()

key, line = reader.read(filename_queue)

x1, x2, y_in = tf.decode_csv(line, record_defaults=[[0.0], [0.0], [0.0]])

sess.run(tf.group(tf.initialize_all_variables(),

                   tf.initialize_local_variables()))

tf.train.start_queue_runners(sess)
```

Reading our data from file (2)

```
... define Lq, trainq ...
```

```
try:
    for i in xrange(100000):
        sess.run(trainq)
        if i % 1000 == 0:
            print "i %d" % i
            eval()
except tf.errors.OutOfRangeError:
    print "done, i=%d" % i
```

Checkpoints

Checkpoints save variable content to disk

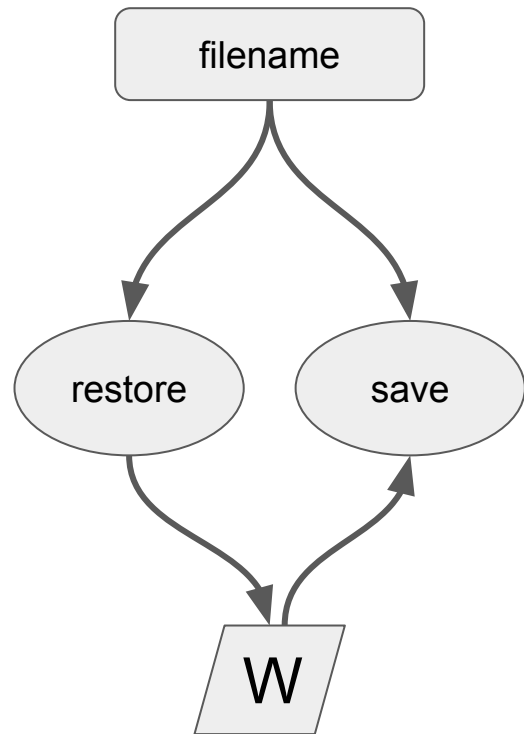
`tf.Saver` adds save/restore Ops to each variable

`Saver.save` feeds a filename and fetches 'save'

`Saver.restore` feeds a filename and fetches 'restore'

You can select which variables to restore

`contrib/learn/utils` has helpers for warmstarting



MNIST

the only benchmark that matters

Load the data

```
mnist = learn.datasets.load_dataset('mnist')
```

```
data = mnist.train.images
```

```
labels = np.asarray(mnist.train.labels, dtype=np.int32)
```

```
eval_data = mnist.test.images
```

```
eval_labels = np.asarray(mnist.test.labels, dtype=np.int32)
```

Linear mnist

```
feature_columns = learn.infer_real_valued_columns_from_input(data)

classifier = learn.LinearClassifier(feature_columns=feature_columns,
                                   n_classes=10)

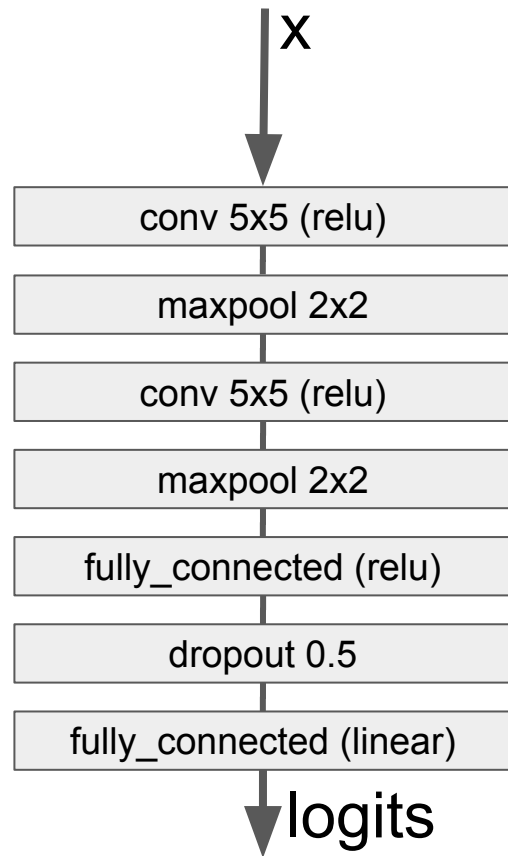
classifier.fit(data, labels, batch_size=100, steps=1000)

classifier.evaluate(eval_data, eval_labels)
```

Let's make a proper model

Use layers to quickly assemble neural networks

```
x = tf.contrib.layers.conv2d(x, kernel_size=[5,5], ...)
x = tf.contrib.layers.max_pool2d(x, kernel_size=[2,2], ...)
x = tf.contrib.layers.conv2d(x, kernel_size=[5,5], ...)
x = tf.contrib.layers.max_pool2d(x, kernel_size=[2,2], ...)
x = tf.contrib.layers.relu(x)
x = tf.contrib.layers.dropout(x, 0.5)
logits = tf.config.layers.linear(x)
```



Use Classifier for training

```
classifier = learn.Classifier(model_fn=model_fn, n_classes=10)
```

model_fn takes (inputs, labels, mode),
returns (logits, loss, train_op)

mode == TRAIN: don't need logits

mode == EVAL: don't need train_op

mode == INFER: only need logits (labels can be None)

How to make loss, train_op

```
def model_fn(inputs, labels, mode)
    logits = <call our logits function>
    onehot = tf.one_hot(labels, 10)
    loss, train_op = None, None

    if mode != tf.contrib.learn.ModeKeys.INFER:
        loss = tf.contrib.losses.softmax_cross_entropy(logits, onehot)

    if mode == tf.contrib.learn.ModeKeys.TRAIN:
        train_op = tf.contrib.layers.optimize_loss(loss,
            tf.contrib.framework.get_global_step(),
            learning_rate=0.001, 'SGD')

    return logits, loss, train_op
```

What I did not tell you

You can control where things are executed:

```
with tf.device('/gpu:2'):
```

You can make a cluster and then distribute computation:

```
with tf.device("/job:worker/task:7"):
```

Using data-parallelism, this is surprisingly automatic

Want more?

Do the DeepDream notebook

Do the tensorflow.org tutorials

Write your next model in TensorFlow